

The LeNet-5 CNN Architecture, developed by Yann LeCun in 1998, is one of the pioneering convolutional neural networks (CNNs).

- The architecture was developed to recognize handwritten and machine-printed characters, a function that showcased the potential of deep learning in practical applications.
- This architecture used backpropagation for training CNNs.
- Suitable for smaller datasets due to limited capacity.

LeNet-5 has a total of 7 layers, excluding the input layer.

- These layers include convolutional layers, pooling (subsampling) layers, and fully connected layers.
- It uses **tanh** activation functions (as opposed to ReLU in modern architectures).

✓ Let us know Layer-by-Layer Description:

1. **Input Layer:** Accepts a 32x32 grayscale image as input. (Images from datasets like MNIST are padded from 28x28 to 32x32 to fit this architecture.)
2. **Convolutional Layer 1 (C1):**
 - Applies 6 filters of size 5x5.
 - Produces an output of 6 feature maps each of size 28x28.
3. **Pooling Layer 1 (S2):**
 - Averages values in 2x2 regions with a stride of 2.
 - Reduces the spatial dimensions from 28x28 to 14x14.
4. **Convolutional Layer 2 (C3):**
 - Applies 16 filters of size 5x5.
 - Produces an output of 16 feature maps each of size 10x10.
5. **Pooling Layer 2 (S4):**
 - Averages values in 2x2 regions with a stride of 2.
 - Reduces the spatial dimensions from 10x10 to 5x5.
6. **Fully Connected Layer 1 (F5):** Flattens the feature maps and passes them through a dense layer with 120 units.
7. **Fully Connected Layer 2 (F6):** Passes the output from the previous layer to another dense layer with 84 units.
8. **Output Layer:**
 - Contains 10 units, representing the 10 digit classes (0–9).
 - Uses a softmax function for classification.

✓ Parameter Sharing

LeNet-5 utilizes shared weights in the convolutional layers to reduce the number of trainable parameters.

To know more:

- <https://www.datasciencecentral.com/lenet-5-a-classic-cnn-architecture/>
- <https://medium.com/@siddheshb008/lenet-5-architecture-explained-3b559cb2d52b>

✓ Image Classification with the CIFAR-10 Dataset Using the LeNet-5 CNN Architecture

```
# Step 1: Importing the necessary libraries
```

```
import tensorflow as tf
from tensorflow.keras import datasets, layers, models
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt
```

```
# Step 2: Loading and pre-processing of the CIFAR-10 dataset
```

```
(train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()
```

```
# Step 3: Normalizing pixel values to be between 0 and 1
```

```
train_images = train_images / 255.0
test_images = test_images / 255.0
```

- One-hot encoding converts categories into a numeric format, enabling models to process them.
- Assigning numeric values (e.g., 1, 2, 3) to categories implies an ordinal relationship that may not exist.
- One-hot encoding removes this unintended relationship by treating each category as independent.
- In classification tasks, the softmax activation function is often used in the output layer.
- One-hot encoding creates a binary vector with a "1" for the true class and "0" elsewhere, which aligns with the softmax output's probability distribution.
- one-hot encoding ensures categorical variables are appropriately represented without introducing unintended relationships or biases, making it an essential preprocessing step for classification tasks.

```
# Step 4: One-hot encode the labels.
```

```
# This step is required to use the loss function "categorical_crossentropy"
```

```
train_labels = to_categorical(train_labels, 10) # here 10 is for- It ensures the output one-hot encoded vectors have a length of 10.
test_labels = to_categorical(test_labels, 10)
```

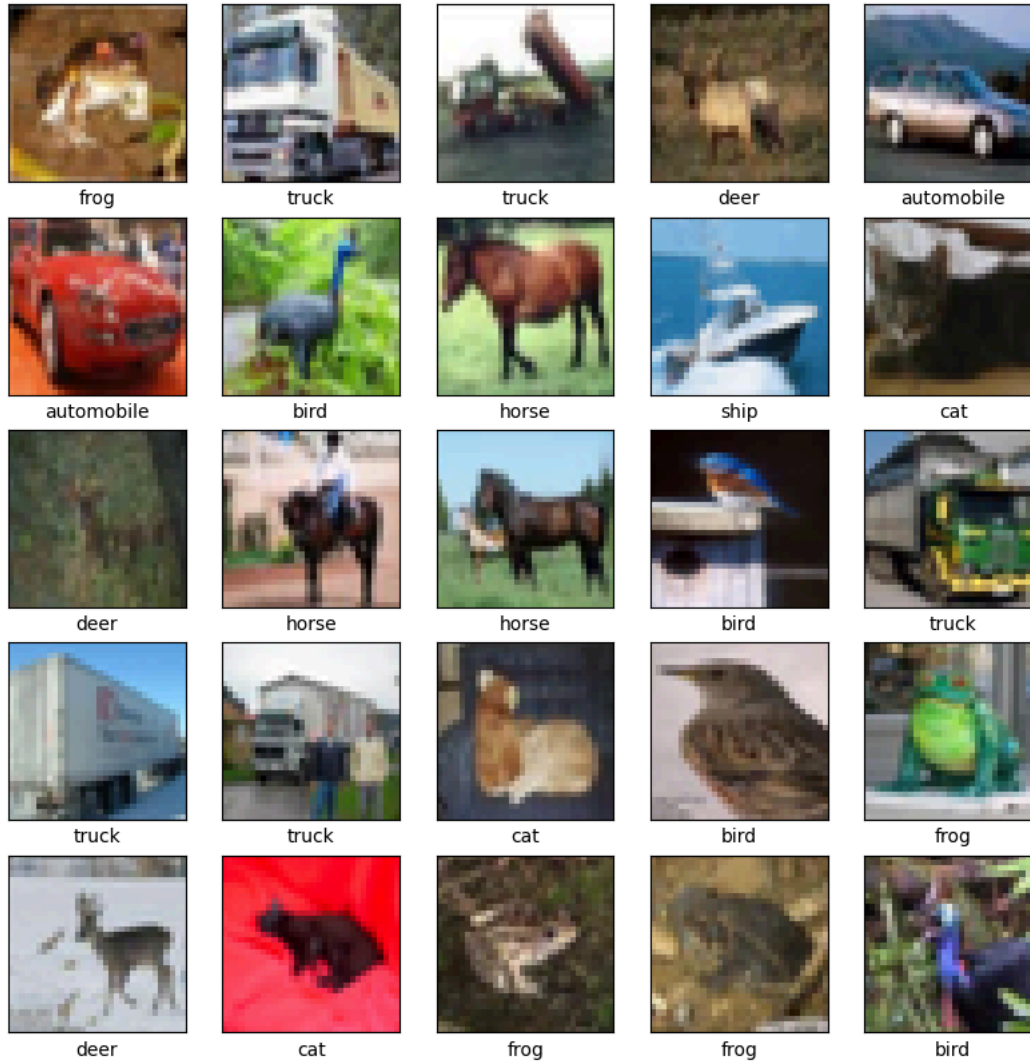
```
# Step 5: Defining the class names for CIFAR-10 images
```

```
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

```
# Step 6: Visualizing a few training images from the CIFAR-10 dataset.

plt.figure(figsize=(10, 10))
for i in range(25):
    plt.subplot(5, 5, i + 1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(train_images[i])
    plt.xlabel(class_names[train_labels[i].argmax()]) # Using argmax to get the label index

plt.show()
```



Step 7: Building the CNN model (LeNet-5 CNN Architecture)

```
model = models.Sequential([
    layers.Conv2D(6, (5, 5), activation='tanh', input_shape=(32, 32, 3)),
    layers.AveragePooling2D((2, 2)),
    layers.Conv2D(16, (5, 5), activation='tanh'),
    layers.AveragePooling2D((2, 2)),
    layers.Conv2D(120, (5, 5), activation='tanh'),
    layers.Flatten(),
    layers.Dense(84, activation='tanh'),
```

```
layers.Dense(10, activation='softmax')
])
```

⚡ /usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential, use `input_shape` argument to the first layer instead.

Step 8: Printing the model summary

```
model.summary()
```

⚡ Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 6)	456
average_pooling2d (AveragePooling2D)	(None, 14, 14, 6)	0
conv2d_1 (Conv2D)	(None, 10, 10, 16)	2,416
average_pooling2d_1 (AveragePooling2D)	(None, 5, 5, 16)	0
conv2d_2 (Conv2D)	(None, 1, 1, 120)	48,120
flatten (Flatten)	(None, 120)	0
dense (Dense)	(None, 84)	10,164
dense_1 (Dense)	(None, 10)	850

Total params: 62,006 (242.21 KB)
 Trainable params: 62,006 (242.21 KB)
 Non-trainable params: 0 (0.00 KB)

Step 9: Compiling the CNN model

```
model.compile(optimizer='adam', # Adam uses a default learning rate of 0.001
loss='categorical_crossentropy',
metrics=['accuracy'])
```

Step 10: Training the CNN model.

```
trained_model = model.fit(train_images, train_labels, epochs=10, batch_size=32, validation_data=(test_images, test_labels))
```

⚡ Epoch 1/10
 1563/1563 ————— 43s 27ms/step - accuracy: 0.3185 - loss: 1.8989 - val_accuracy: 0.4045 - val_loss: 1.6714
 Epoch 2/10
 1563/1563 ————— 79s 25ms/step - accuracy: 0.4276 - loss: 1.6227 - val_accuracy: 0.4657 - val_loss: 1.4975
 Epoch 3/10
 1563/1563 ————— 41s 25ms/step - accuracy: 0.4845 - loss: 1.4483 - val_accuracy: 0.4910 - val_loss: 1.4253
 Epoch 4/10
 1563/1563 ————— 40s 24ms/step - accuracy: 0.5204 - loss: 1.3544 - val_accuracy: 0.5173 - val_loss: 1.3641
 Epoch 5/10
 1563/1563 ————— 38s 25ms/step - accuracy: 0.5399 - loss: 1.2963 - val_accuracy: 0.5255 - val_loss: 1.3505
 Epoch 6/10
 1563/1563 ————— 41s 25ms/step - accuracy: 0.5633 - loss: 1.2382 - val_accuracy: 0.5232 - val_loss: 1.3531
 Epoch 7/10

```

1563/1563 ————— 41s 25ms/step - accuracy: 0.5741 - loss: 1.1992 - val_accuracy: 0.5288 - val_loss: 1.3335
Epoch 8/10
1563/1563 ————— 41s 25ms/step - accuracy: 0.5936 - loss: 1.1511 - val_accuracy: 0.5311 - val_loss: 1.3363
Epoch 9/10
1563/1563 ————— 40s 26ms/step - accuracy: 0.6034 - loss: 1.1187 - val_accuracy: 0.5293 - val_loss: 1.3479
Epoch 10/10
1563/1563 ————— 40s 25ms/step - accuracy: 0.6174 - loss: 1.0918 - val_accuracy: 0.5461 - val_loss: 1.3261

```

```
# Step 11: Evaluating the performance of the CNN model
```

```
test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
print(f'\nTest accuracy is: {test_acc}')
```

```
313/313 - 3s - 9ms/step - accuracy: 0.5461 - loss: 1.3261
```

```
Test accuracy is: 0.5461000204086304
```

```
# Plotting training and validation accuracy
```

```
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(trained_model.history['accuracy'], label='Training Accuracy')
plt.plot(trained_model.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.ylim([0, 1])
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')
plt.grid(True)
```

```
# Plotting training and validation loss
```

```
plt.subplot(1, 2, 2)
plt.plot(trained_model.history['loss'], label='Training Loss')
plt.plot(trained_model.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.ylim([0, 2])
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.grid(True)
plt.show()
```